

# PARALLDROID

Towards a unified heterogeneous development model in Android™

---

Alejandro Acosta  
aacostad@ull.es


Francisco Almeida  
falmeida@ull.es



Universidad  
de La Laguna

High Performance Computing Group

# Where we come from



ULL

# Who we are

High Performance Computing Group



# Outline


- Motivation
- Android Programming Models
- Paralldroid
  - Directives
- Extending to Classes
  - Directives
- Computational Results
- Future

# MOTIVATION

---

# The Goal: To ease the Parallelization in Mobile Devices

- Multicore + GPU



# The Context: Android Devices

- Sequential programmers (no knowledge about parallelism)
- Applications demanding more computational capabilities
  - Image/Video Processing
  - Augmented Reality
  - ...
- Different sequential and parallel programming models
- Difficulties - Programmability Wall
  - Developing new efficient code is a difficult task
  - Adapting existent code to new emergent architectures is also difficult

# The Context: Android Devices

- Sequential programmers (no knowledge about parallelism)
- Applications demanding more computational capabilities
  - Image/Video Processing
  - Augmented Reality
  - ...
- Different sequential and parallel programming models
- Difficulties - Programmability Wall
  - Developing new efficient code is a difficult task
  - Adapting existent code to new emergent architectures is also difficult
- The same scenario than in traditional scientific applications???



# The Hypothesis: To Apply the Known Methodologies

- Scientific Context:
  - Many tools developed to ease the programmer task
  - Standards (based in compiler directives) designed to simplify parallel programming
    - OpenMP: Shared memory systems
    - OpenACC: Accelerator systems

# The Hypothesis: To Apply the Known Methodologies

- Scientific Context:
  - Many tools developed to ease the programmer task
  - Standards (based in compiler directives) designed to simplify parallel programming
    - OpenMP: Shared memory systems
    - OpenACC: Accelerator systems
- To extend these ideas to the Android programming models under a unified framework

# THE ANDROID PROGRAMMING MODELS

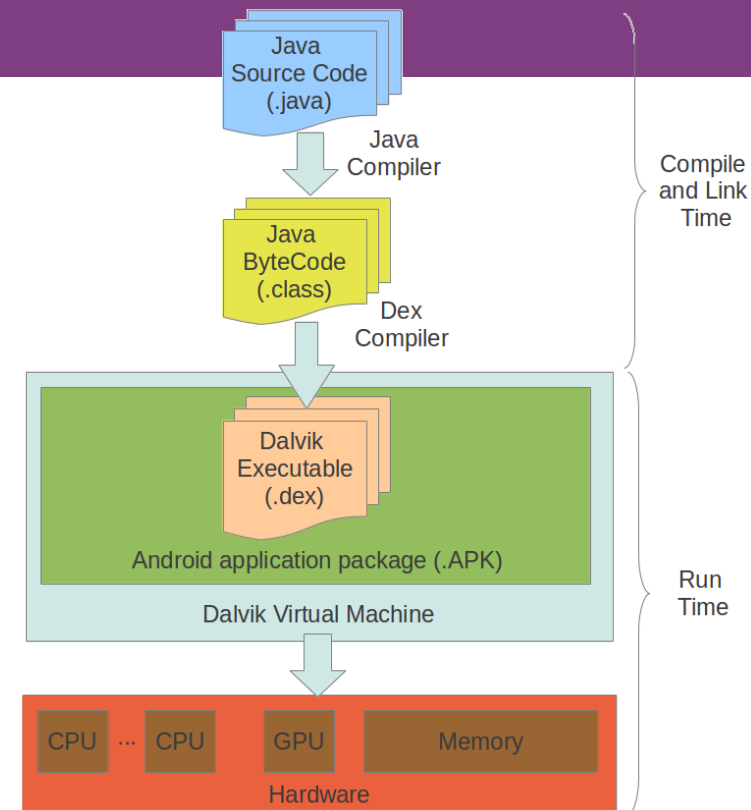
---

# Android Programming Models

- Java (Dalvik)
- Native C
- Renderscript

# Java

- Object Oriented
- Well-known.
- Rapid learning curve.
- Large developers community.
- Main programming language for Android.



Java code

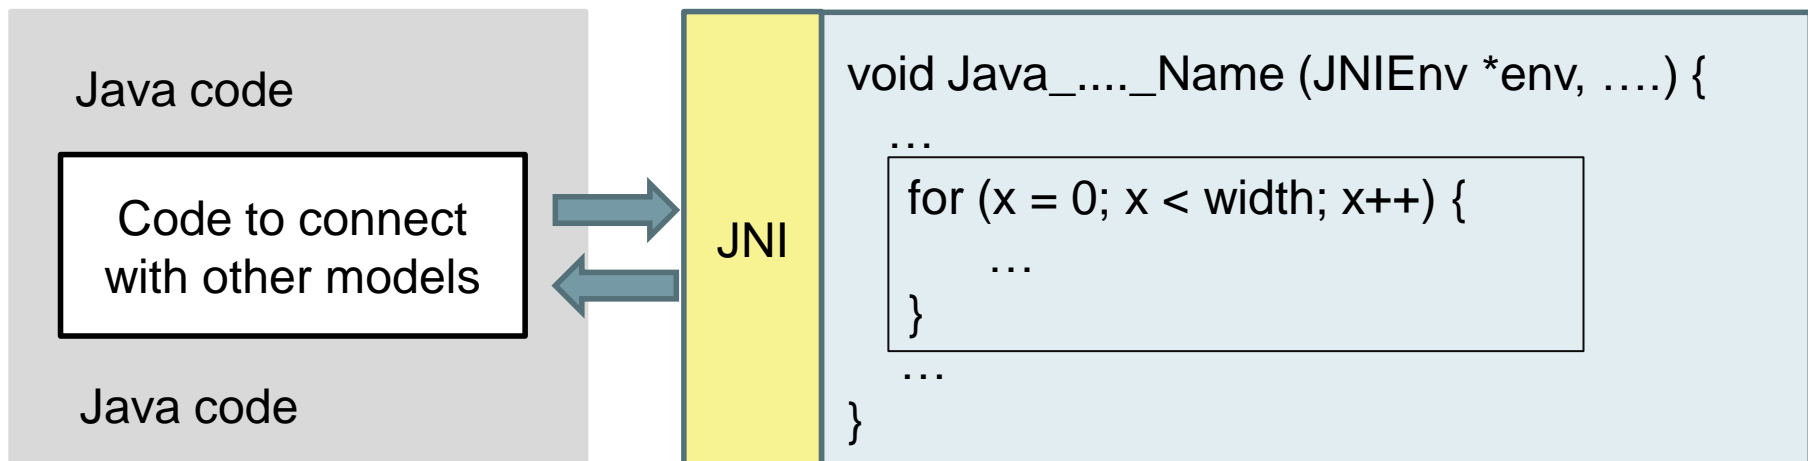
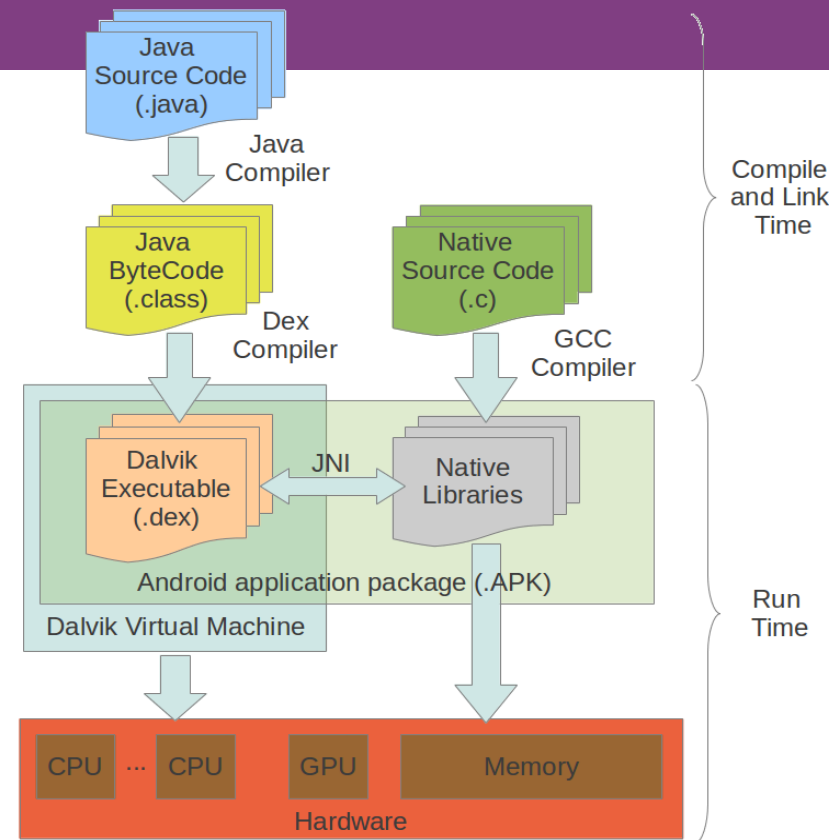
```
for (x = 0; x < width; x++) {  
    for (y = 0; y < height; y++) {  
        ...  
    }  
}
```

Java code

ULL

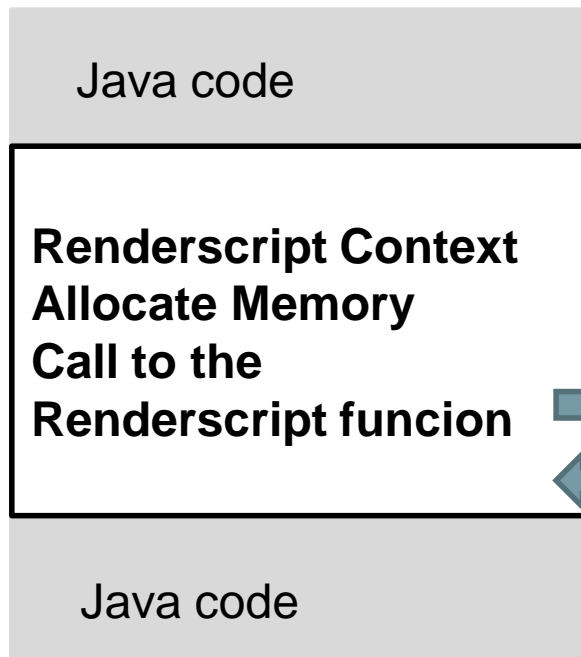
# Native C

- Well-known.
- Large developers community.
- Used to implement some sections of the application.
- Use JNI.
- Interfaces to access the Java Objects
- Compatibility with C libraries
- OpenCL (Blocked on Nexus with Android 4.3)

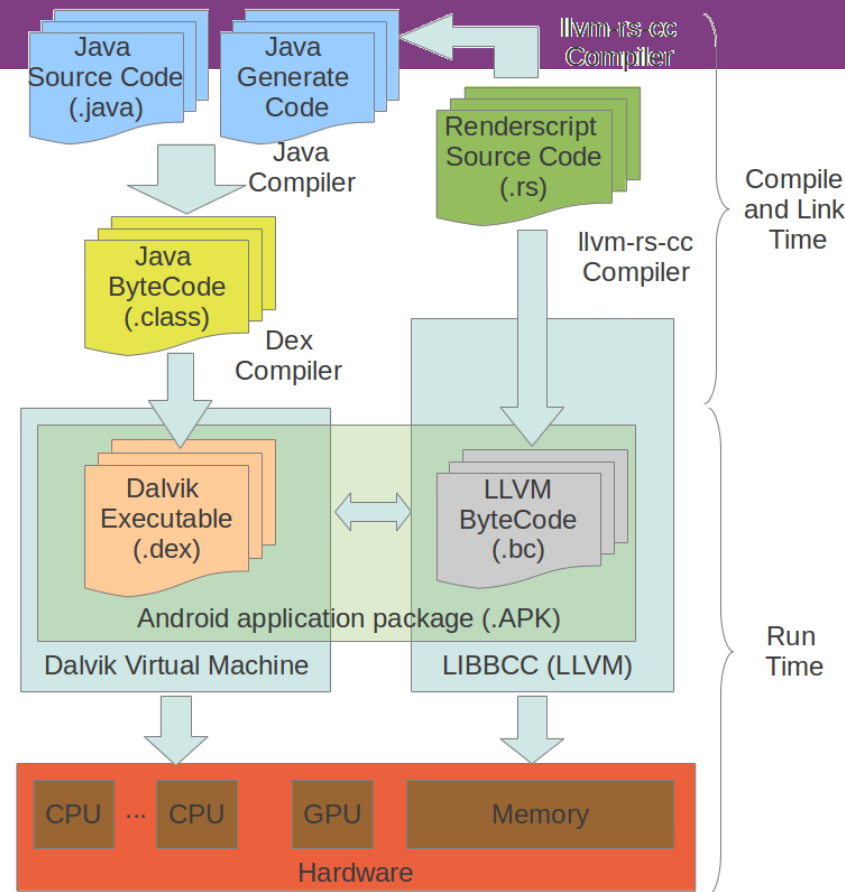


# RenderScript

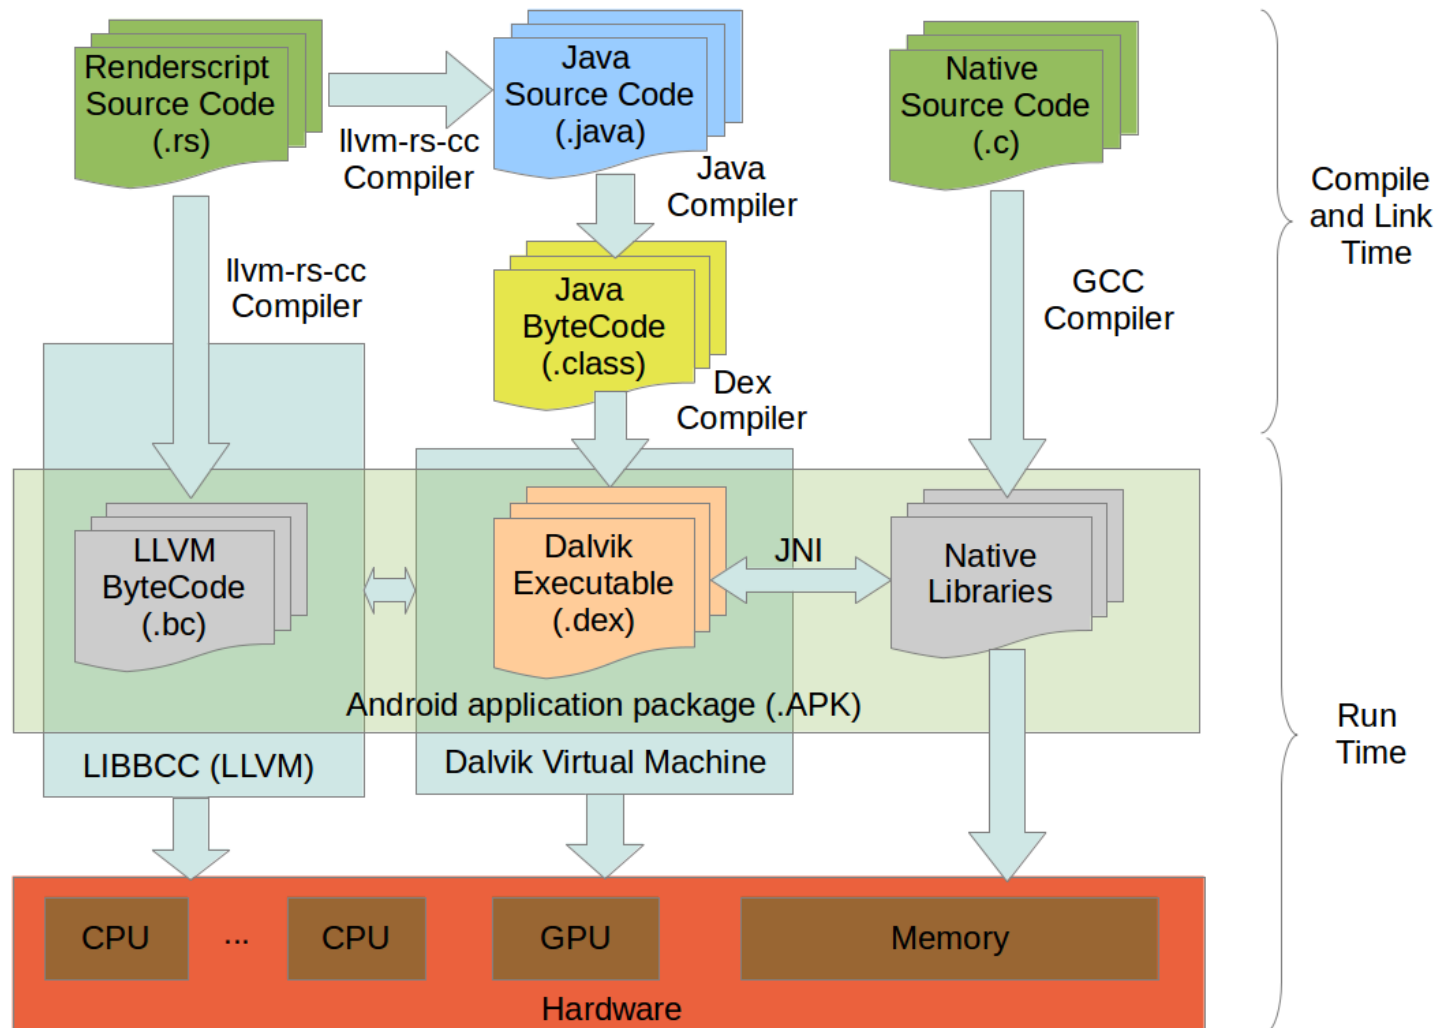
- High performance.
- Similar to C with some extension.
- Support for parallel executions.
- Support for GPU executions.
- Some Java Objects are ported to the RenderScript layer



```
void root(const uchar4 *v_in, uchar4 *v_out) {  
    ...  
    for (int x = 0; x < width; x++) {  
        ...  
    }  
    ...  
}
```

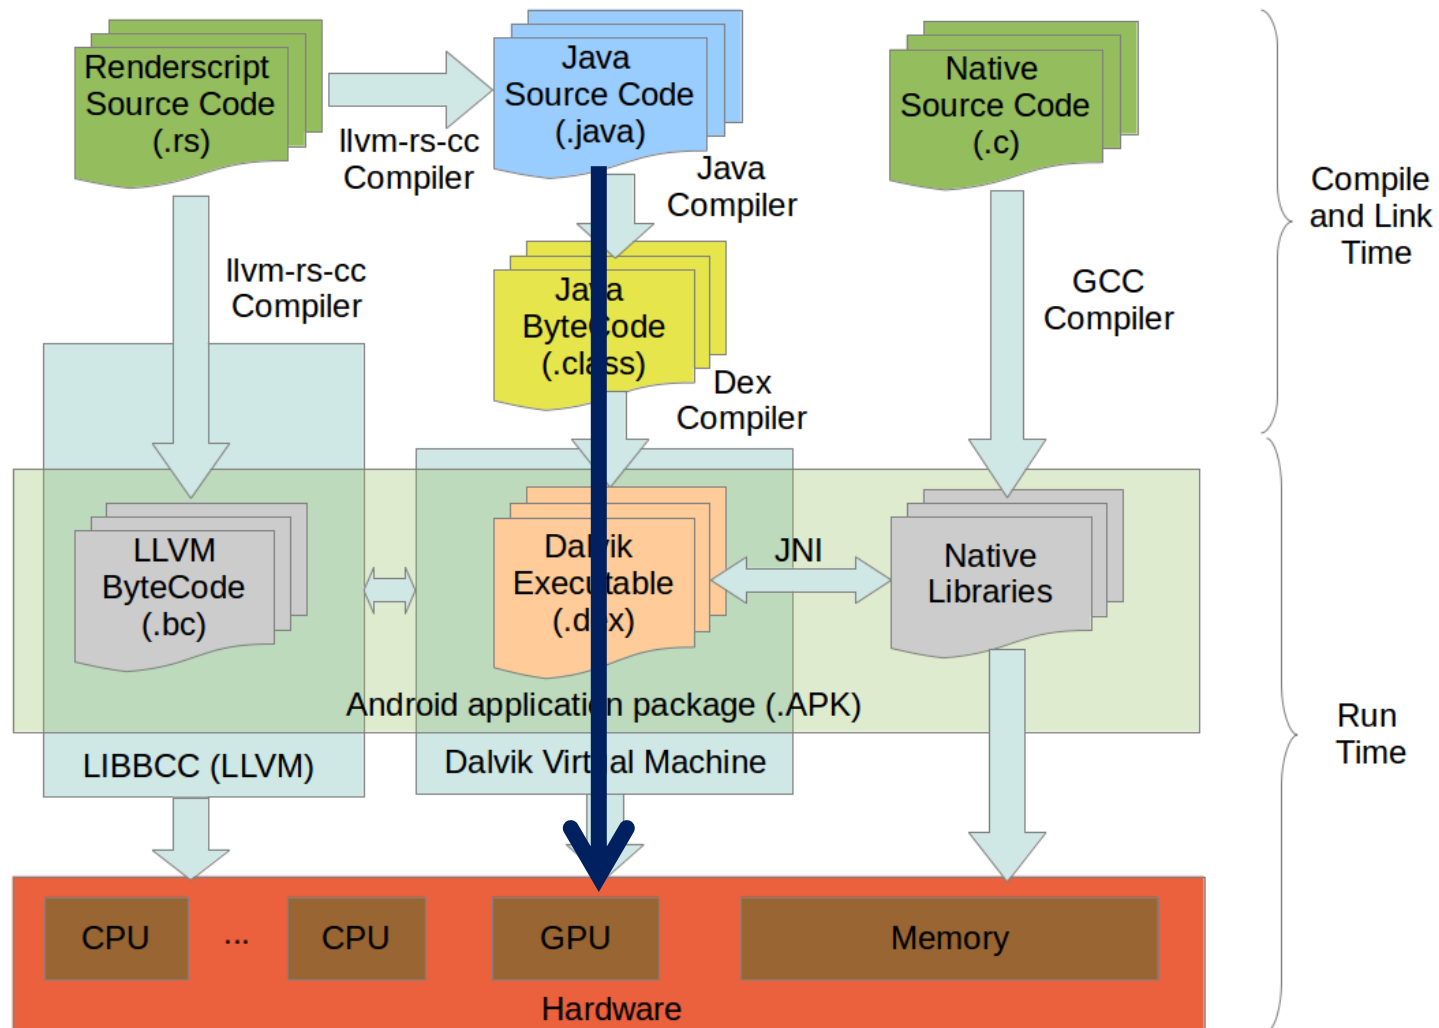


# Compilation Process

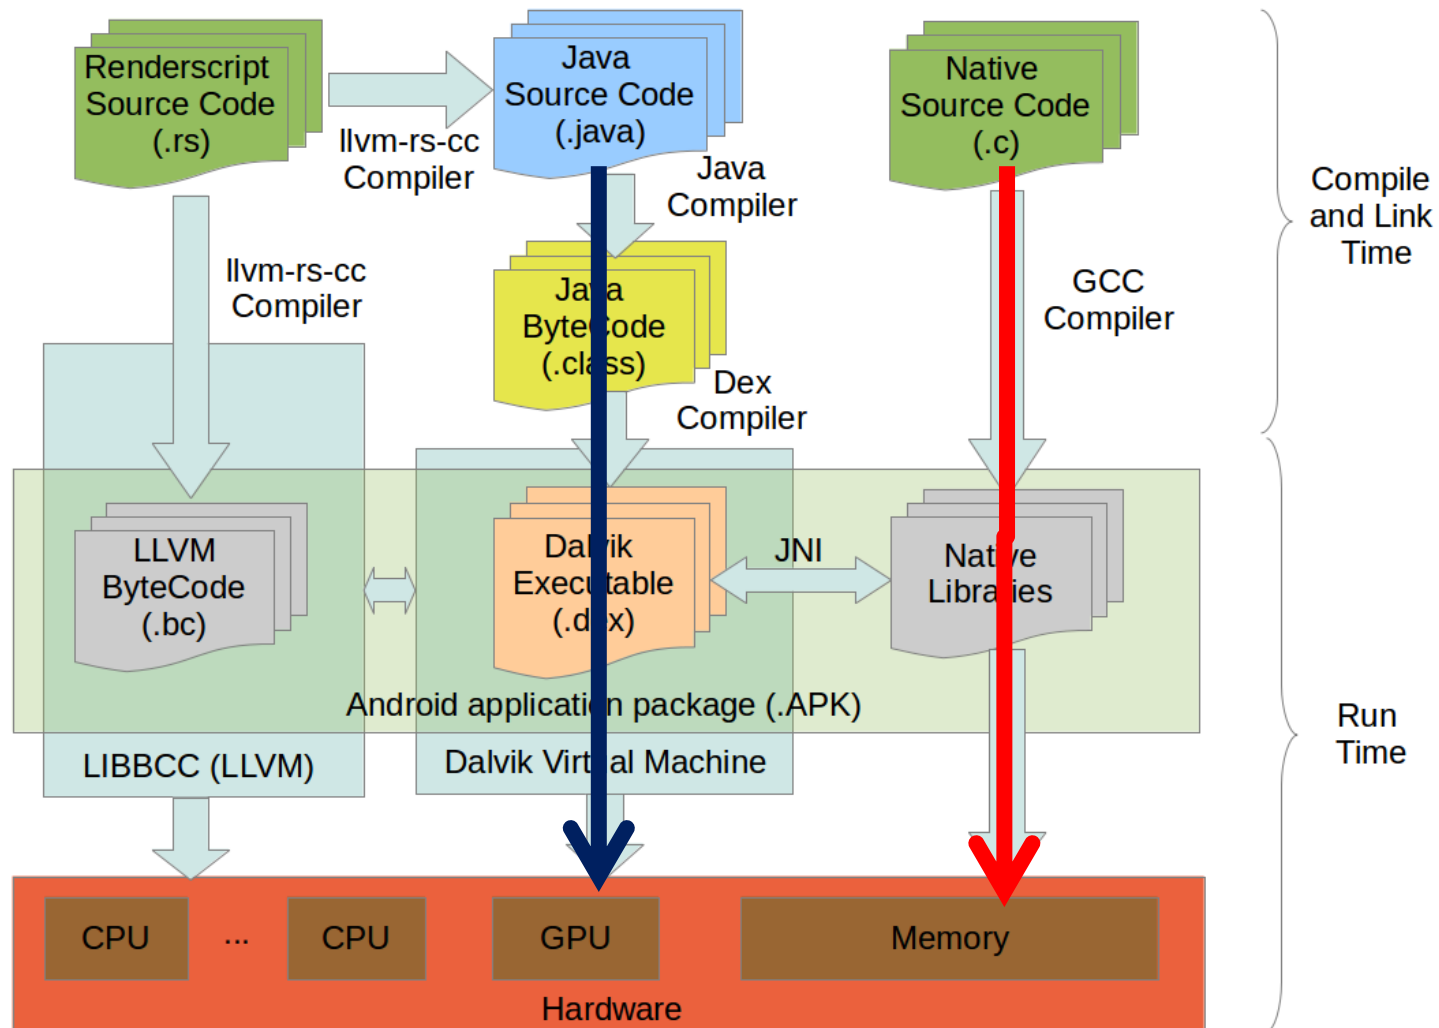




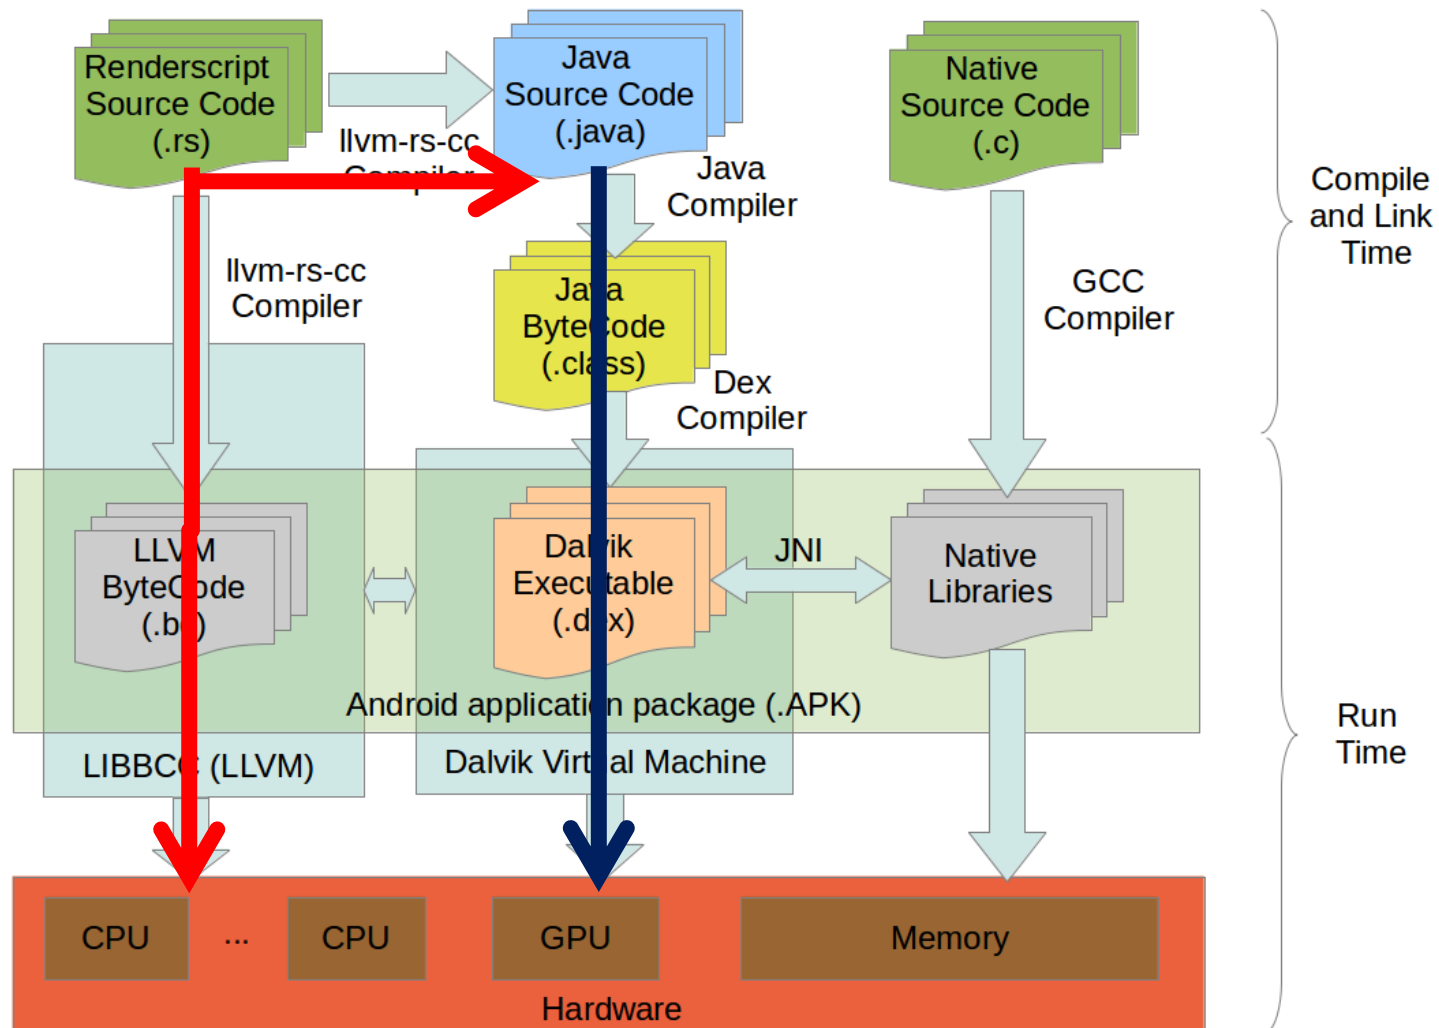
# Compilation Process



# Compilation Process



# Compilation Process



# The Test: Nexus 7

- Renderscript ImageProcessing benchmark

(AOSP: frameworks/base/tests/RenderScriptTests/ImageProcessing)

- Grayscale
- Convolve 3x3
- Convolve 5x5
- Levels

Java (Dalvik)  
Native C  
Renderscript

- General Convolve

- 3x3
- 5x5
- 7x7
- 9x9


**Gray scale**




**1600x1067**

ULL

## AOSP Benchmark



## General convolve



■ Native C


■ Renderscript Sequential

■ Renderscript Parallel

Programming models	CPU 1	CPU 2	CPU 3	CPU 4
Java	100%	Offline	Offline	Offline
Native C	100%	Offline	Offline	Offline
Renderscript sequential	100%	Offline	Offline	Offline
Renderscript Parallel	100%	100%	100%	100%



# Implementation



# PARALLDROID

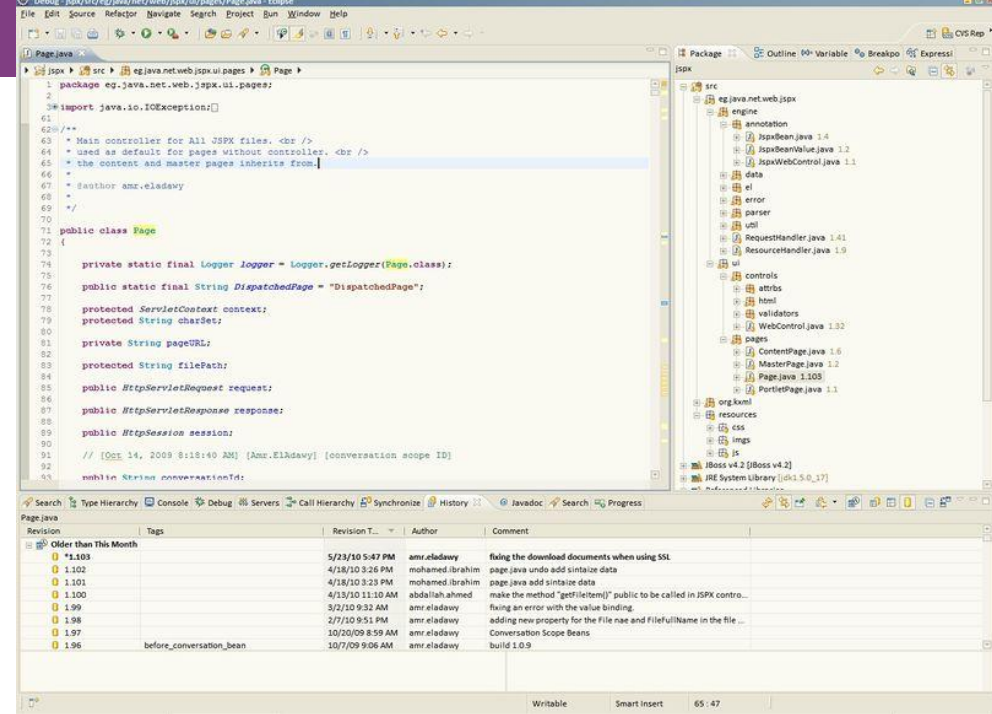
---

*Towards a unified heterogeneous development model in Android.*  
HeteroPar 2013

*Paralldroid: A Framework for Parallelism in Android.* LEAP 2013.  
(Low Energy Application Parallelism)

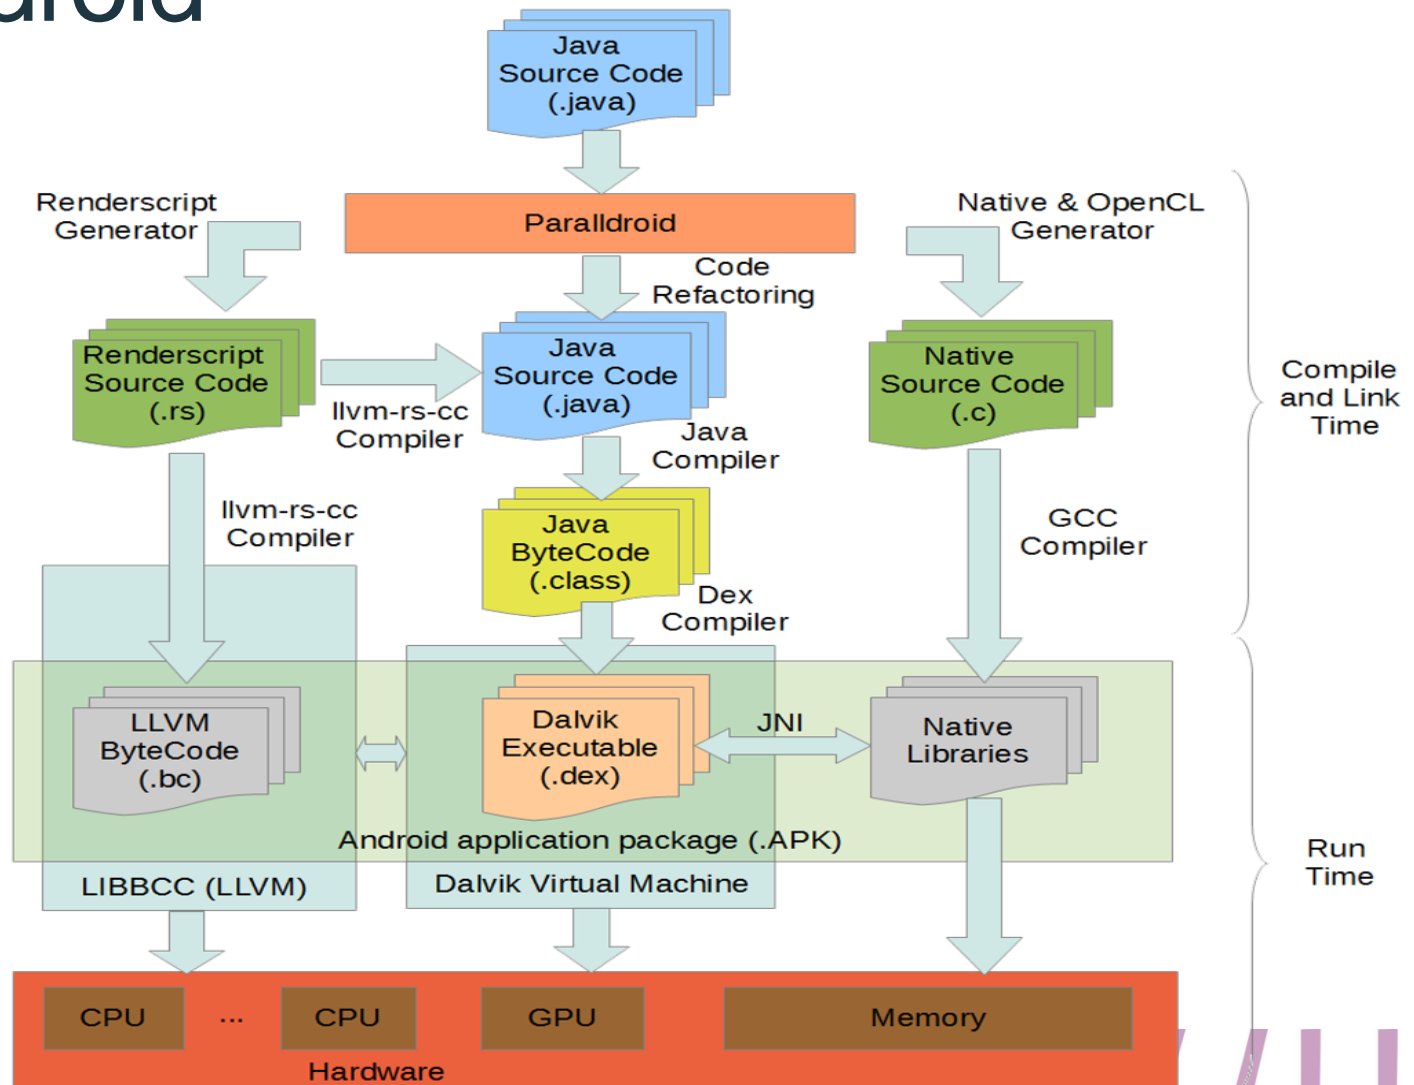
# Paralldroid

- Framework based on Java.
- Source to Source translator.
- Eclipse plugin.
- Annotated Java code.
- OpenMP 4.0 extension
- **Advantages:**
  - Increased use of the parallel devices by non-expert users.
  - Rapid inclusion of emerging technology into their systems.
  - Delivery of new applications due to the rapid development time.
  - Unify the different programming models of Android.
- **Disadvantages:**
  - Less performance compared with an adhoc version (at a low effort).
  - Eclipse dependency.

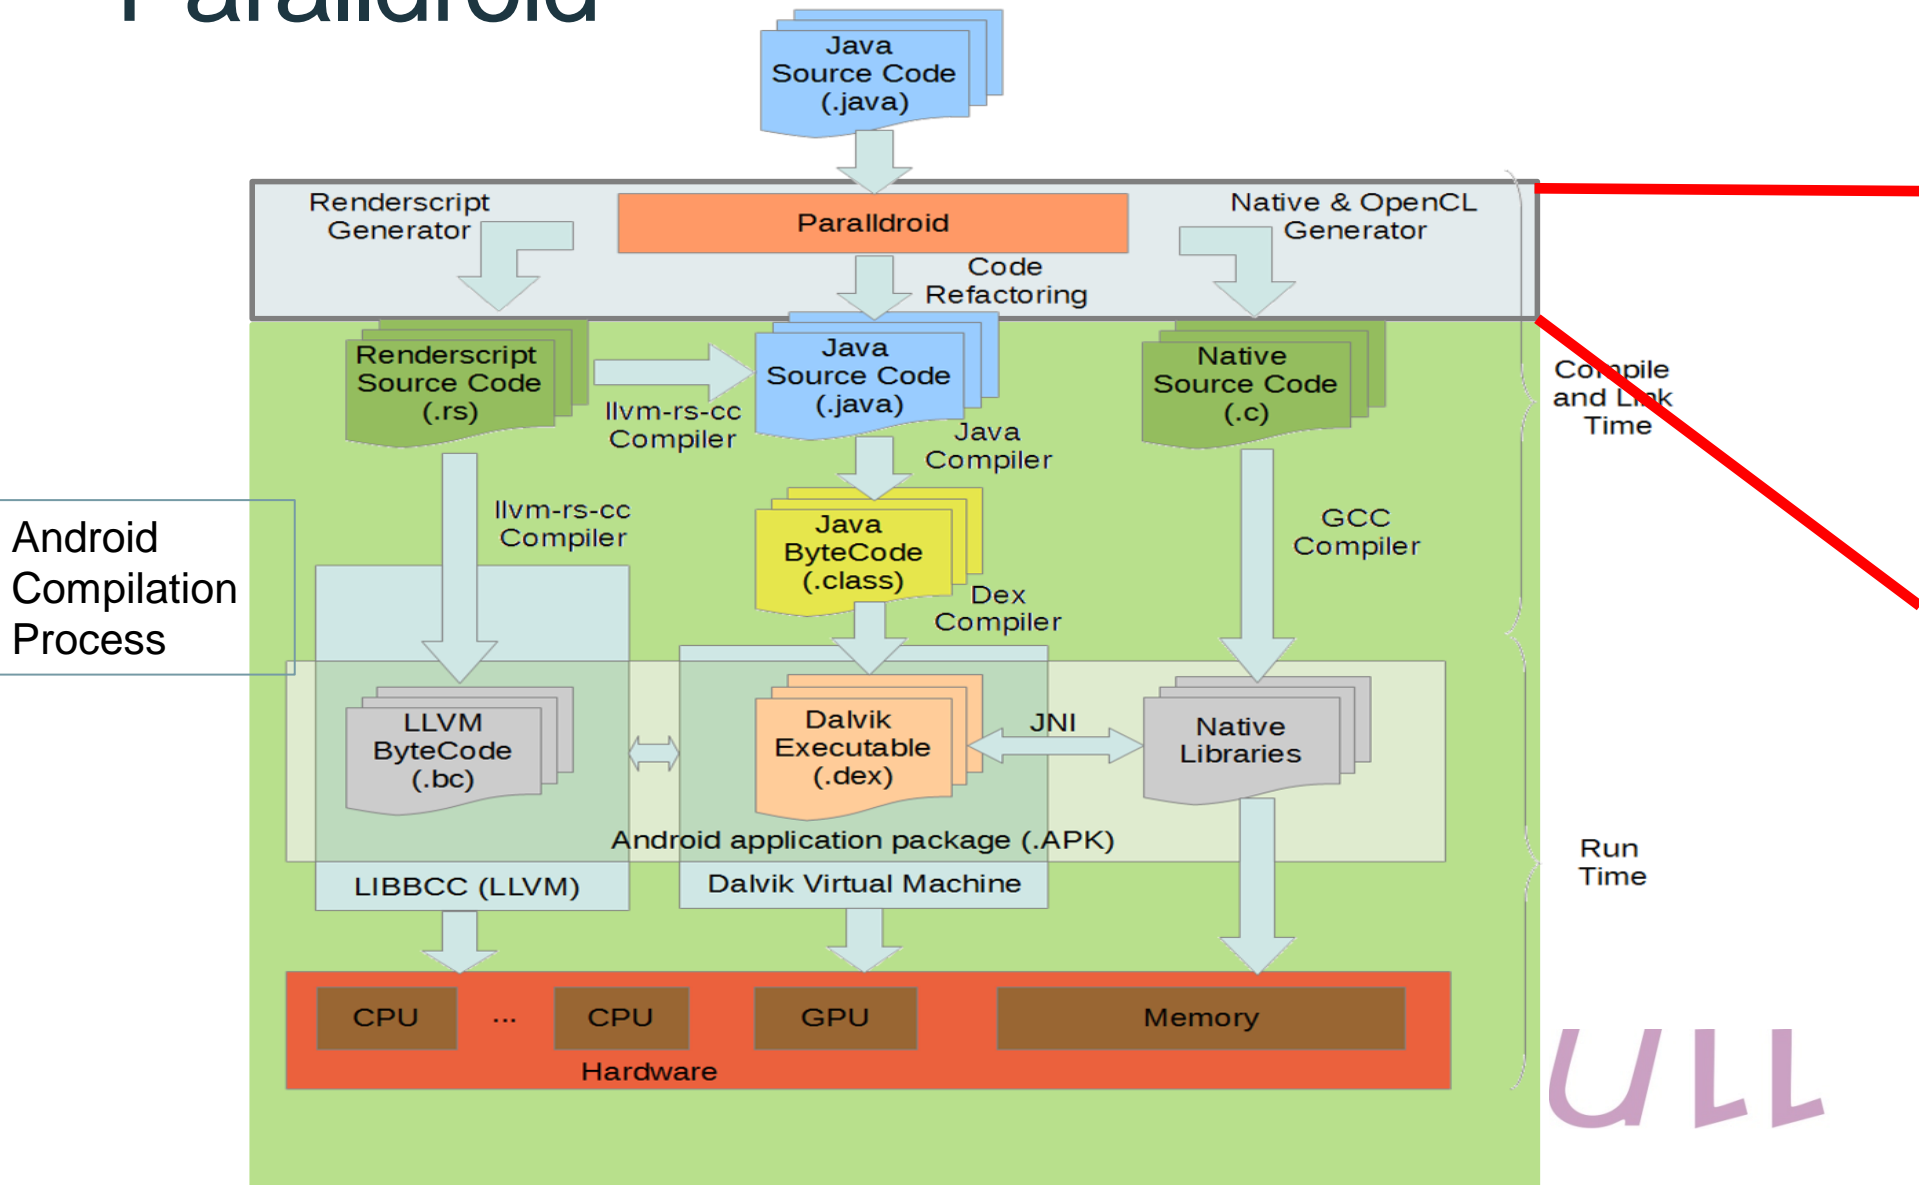




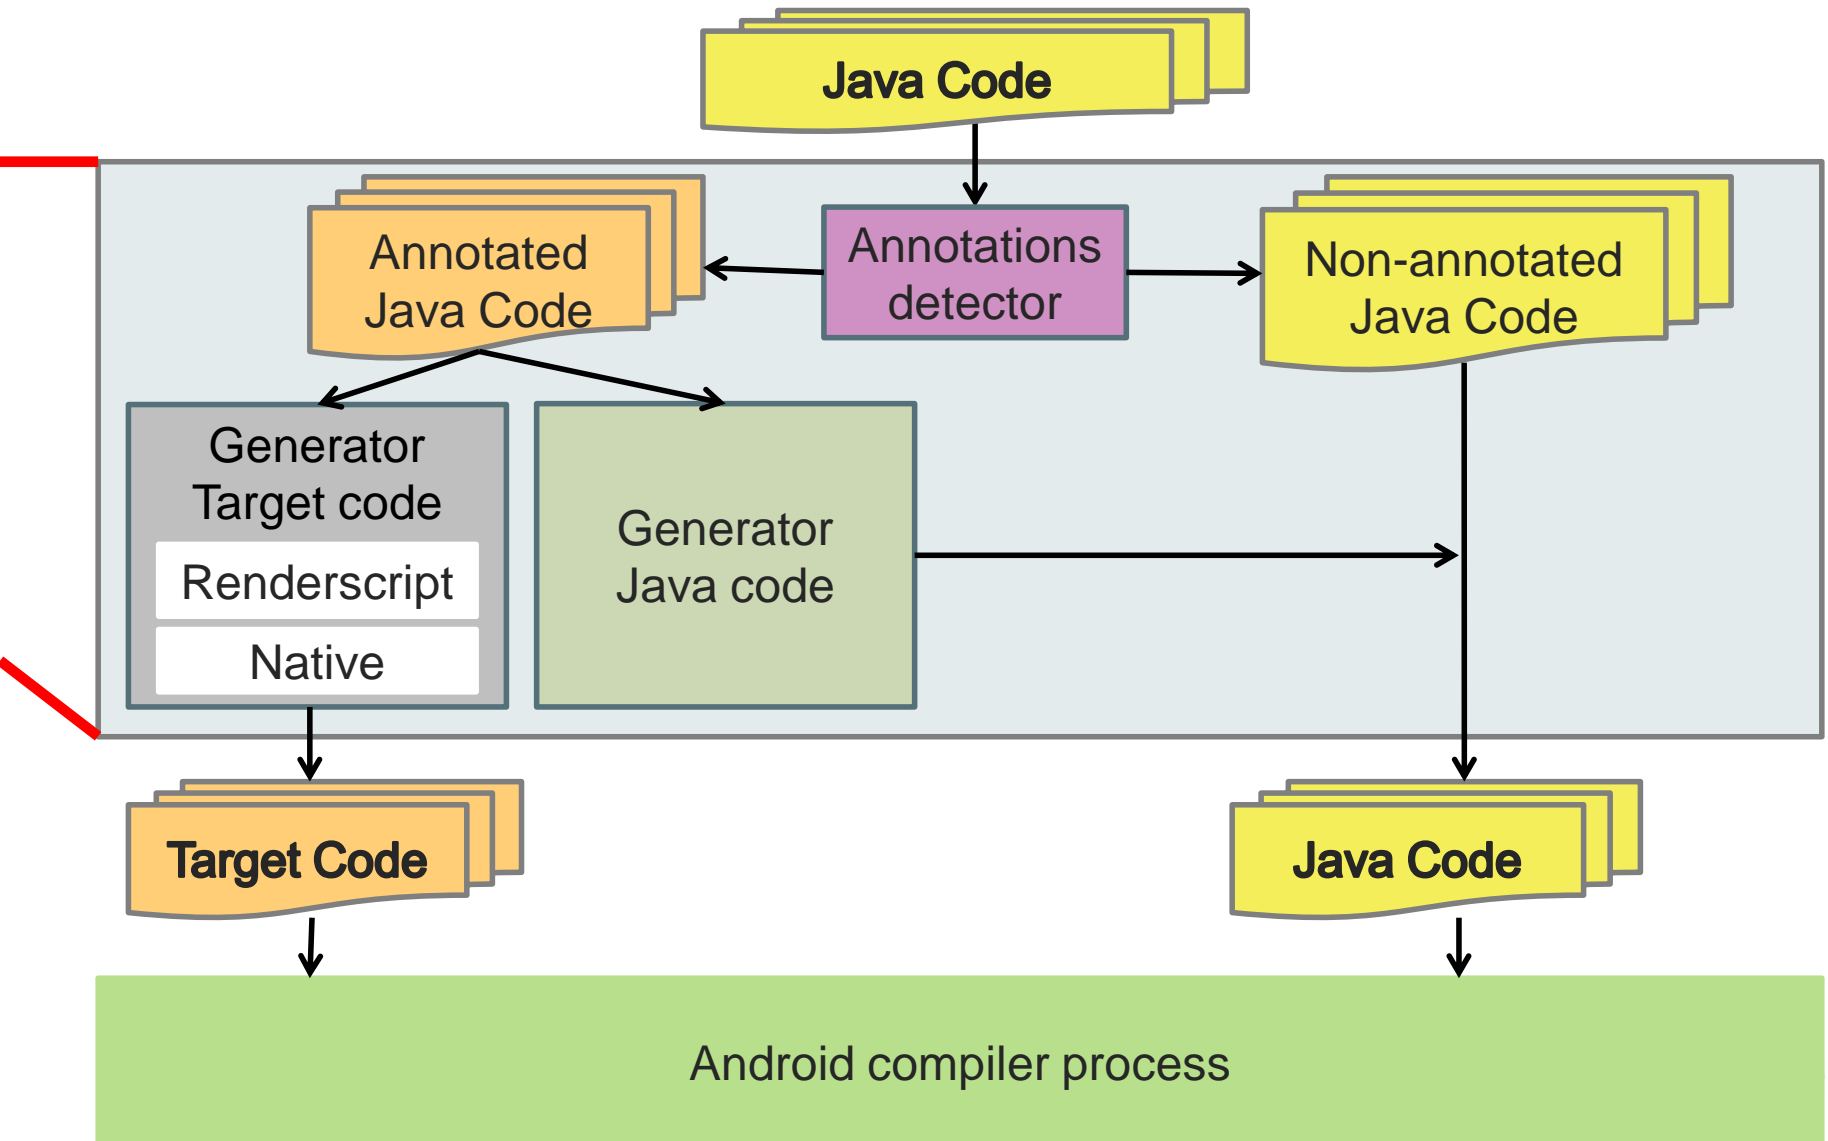
# Paralldroid



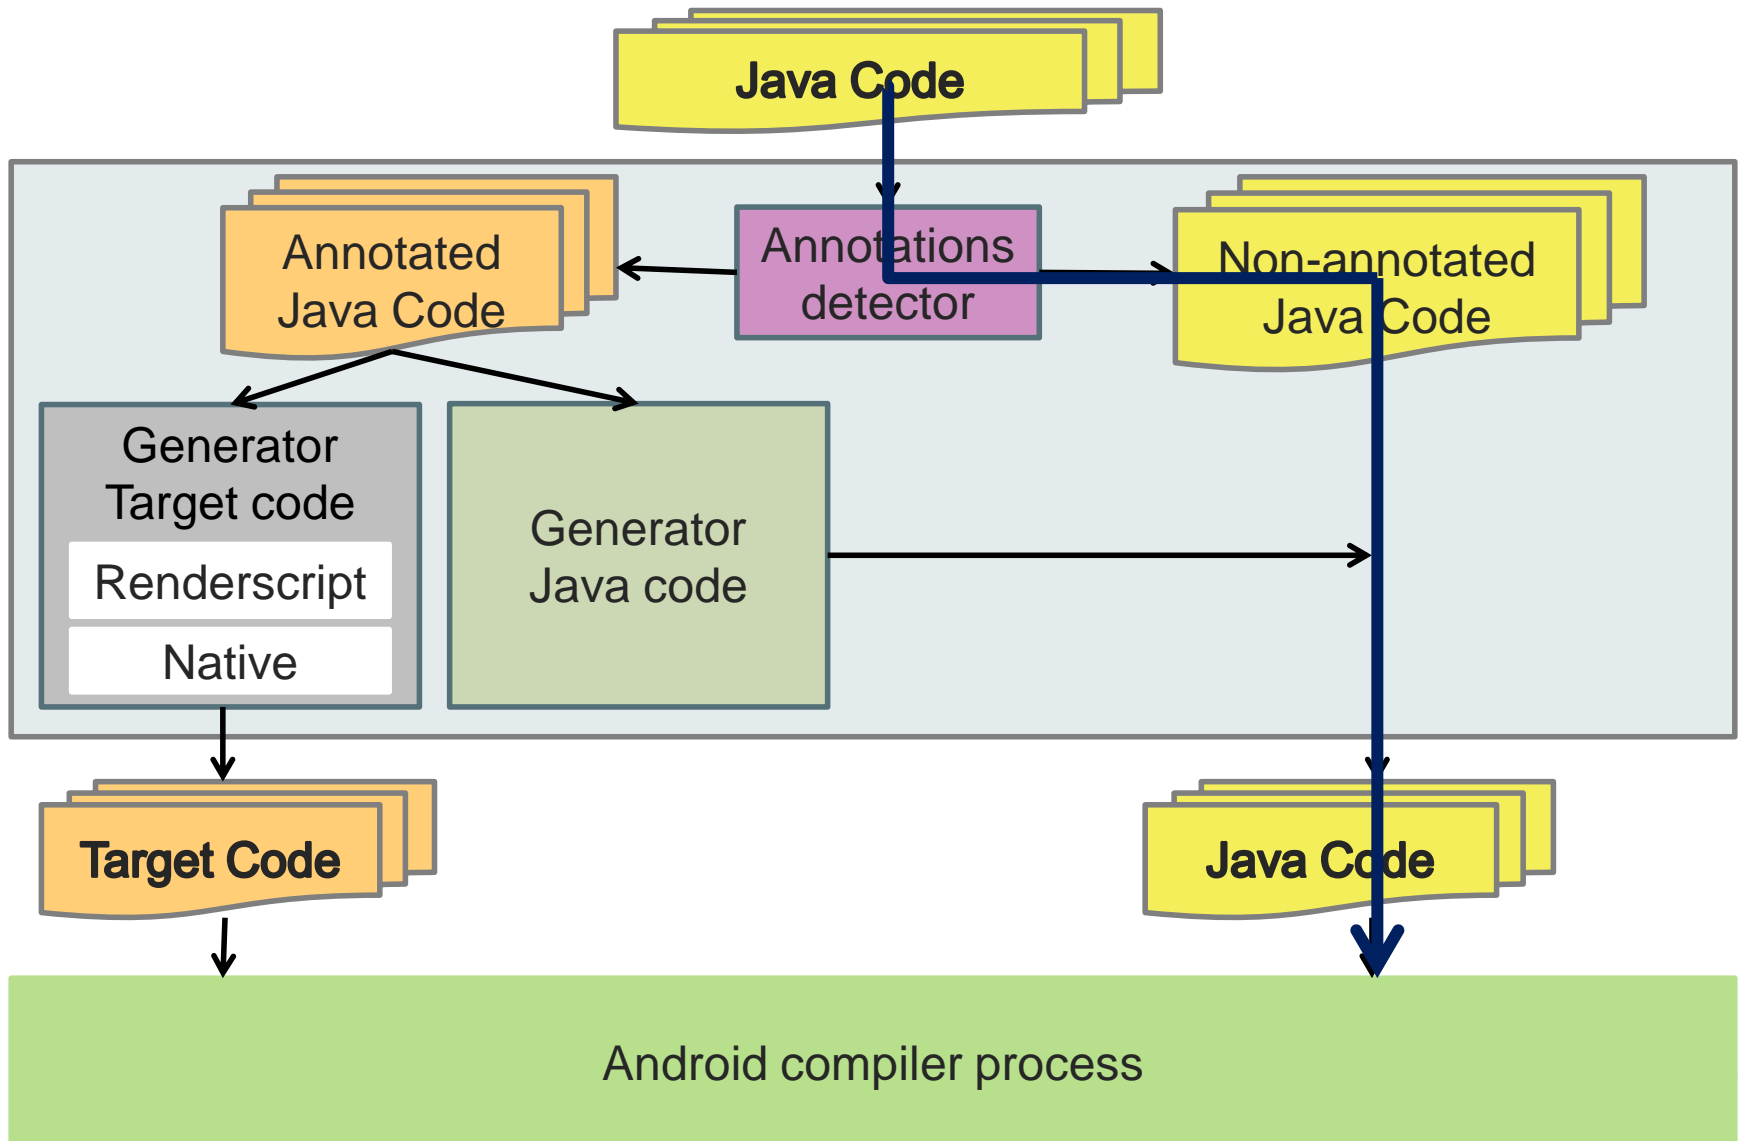
# Paralldroid



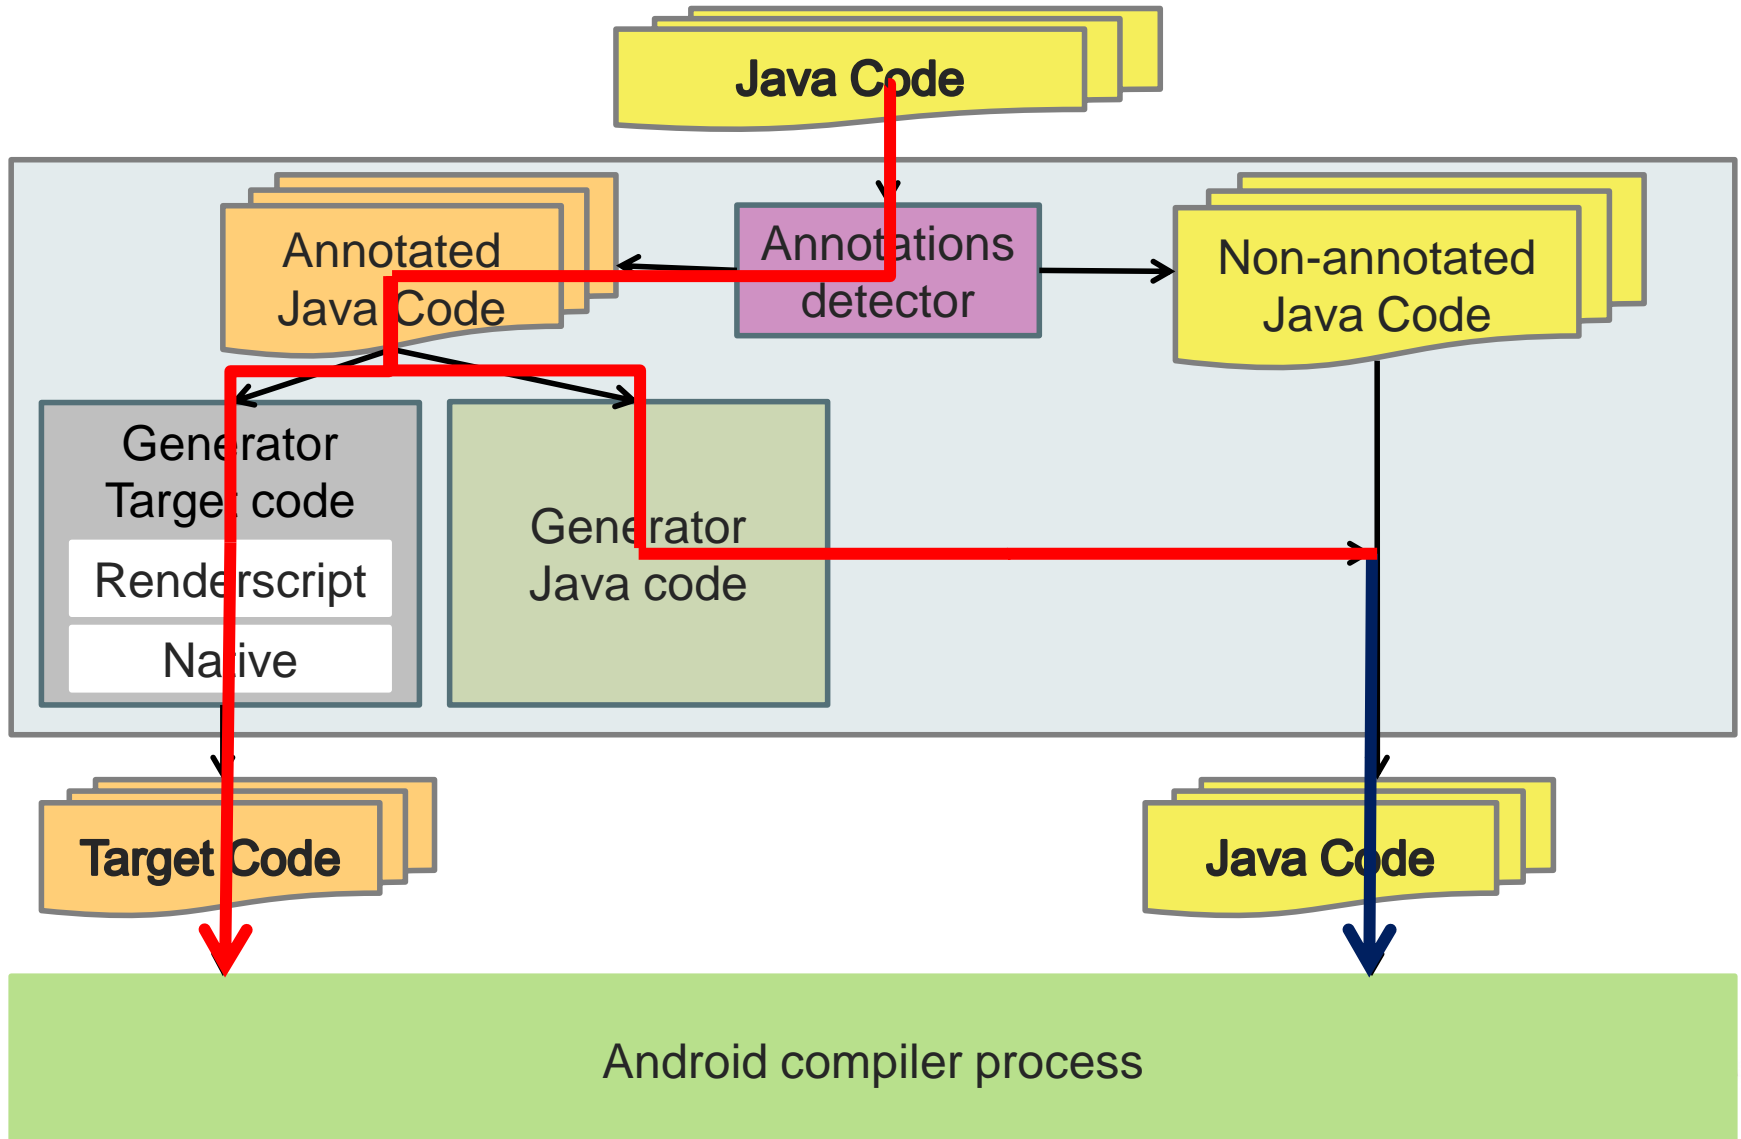
# Paralldroid compiler process



# Non-annotated Java code




# Annotated Java code



# Directive: Target data

- Create a data environment.
- Mapping data to the target context.
- Clauses
  - **Lang**
    - Extension to the OpenMP standard,
    - Target language (Renderscript, Native or OpenCL)
  - Map
    - Maps a variable from the current Java context to the target data context.
    - Map types: Alloc, To, From, ToFrom (default).



# Directive: Target


- Create a data environment.
- Mapping data to the target context.
- Execute the code associate to the directive
- Clauses
  - **Lang**
    - Extension to the OpenMP standard,
    - Target language (Renderscript, Native or OpenCL)
  - Map
    - Maps a variable from the current Java context to the target data context.
    - Map types: Alloc, To, From, ToFrom (default).



ULL

# Directive: Parallel for


- Used in the context of a target directive
- Distributing the load of the for loop between the threads available
- Clauses
  - Private
  - Firstprivate
  - Shared
  - Collapse
  - **Rsvector**
    - Extension to the OpenMP standard.
    - Input and output vectors used in Renderscript.






# Directive: Teams

- Used in the context of a target directive
- Create teams or groups of threads.
- Clauses
  - Num\_teams
  - Num\_thread
  - Private
  - Firstprivate
  - Shared.



# Directive: Distribute

- Used in the context of a teams directive
- Distributing the load of the for loop between the teams available
- Clauses
  - Private
  - Firstprivate
  - Colapse



# Grayscale Java implementation

```
public void grayscale() {  
    int pixel, sum, x;  
    int [] scrPxs = new int[width*height];  
    int [] outPxs = new int[width*height];  
    bitmapIn.getPixels(scrPxs, 0, width, 0, 0, width, height);  
  
    for(x = 0; x < width*height; x++) {  
        pixel = scrPxs[x];  
        sum = (int)(((pixel) & 0xff) * 0.299f);  
        sum += (int)(((pixel >> 8) & 0xff) * 0.587f);  
        sum += (int)(((pixel >> 16) & 0xff) * 0.114f);  
        outPxs[x] = (sum) + (sum << 8) + (sum << 16) + (scrPxs[x] & 0xff000000);  
    }  
  
    bitmapOut.setPixels(outPxs, 0, width, 0, 0, width, height);  
}
```

# Grayscale Java implementation

```
public void grayscale() {  
    int pixel, sum, x;  
    int [] scrPxs = new int[width*height];  
    int [] outPxs = new int[width*height];  
    bitmapIn.getPixels(scrPxs, 0, width, 0, 0, width, height);  
  
    for(x = 0; x < width*height; x++) {  
        pixel = scrPxs[x];  
        sum = (int)(((pixel) & 0xff) * 0.299f);  
        sum += (int)(((pixel >> 8) & 0xff) * 0.587f);  
        sum += (int)(((pixel >> 16) & 0xff) * 0.114f);  
        outPxs[x] = (sum) + (sum << 8) + (sum << 16) + (scrPxs[x] & 0xff000000);  
    }  
  
    bitmapOut.setPixels(outPxs, 0, width, 0, 0, width, height);  
}
```

# Native

```
public void grayscale() {  
    int pixel, sum, x;  
    int [] scrPxs = new int[width*height];  
    int [] outPxs = new int[width*height];  
    bitmapIn.getPixels(scrPxs, 0, width, 0, 0, width, height);  
  
    // pragma paraldroid target lang(native) map(alloc:x,pixel,sum)  
    for(x = 0; x < width*height; x++) {  
        pixel = scrPxs[x];  
        sum = (int)((((pixel) & 0xff) * 0.299f);  
        sum += (int)((((pixel >> 8) & 0xff) * 0.587f);  
        sum += (int)((((pixel >> 16) & 0xff) * 0.114f);  
        outPxs[x] = (sum) + (sum << 8) + (sum << 16) + (scrPxs[x] & 0xff000000);  
    }  
  
    bitmapOut.setPixels(outPxs, 0, width, 0, 0, width, height);  
}
```

# RenderScript

```
public void grayscale() {  
    int pixel, sum, x;  
    int [] scrPxs = new int[width*height];  
    int [] outPxs = new int[width*height];  
    bitmapIn.getPixels(scrPxs, 0, width, 0, 0, width, height);  
  
    // pragma paraldroid target lang(rs) map(to:scrPxs,width,height) map(from:outPxs)  
    // pragma paraldroid parallel for private(x,pixel,sum) rsvector(scrPxs,outPxs)  
    for(x = 0; x < width*height; x++) {  
        pixel = scrPxs[x];  
        sum = (int)((((pixel) & 0xff) * 0.299f);  
        sum += (int)((((pixel >> 8) & 0xff) * 0.587f);  
        sum += (int)((((pixel >> 16) & 0xff) * 0.114f);  
        outPxs[x] = (sum) + (sum << 8) + (sum << 16) + (scrPxs[x] & 0xff000000);  
    }  
  
    bitmapOut.setPixels(outPxs, 0, width, 0, 0, width, height);  
}
```



# ONE STEP FORWARD: EXTENDING TO CLASSES

---

[illegible]

- # Android Studio

ULL



# Grayscale Java implementation

```
public class GrayScale {  
    private final float gMonoMult[] = {0.299f, 0.587f, 0.114f};  
    private int width;  
    private int height;  
  
    public GrayScale (int width, int height) {  
        this.width = width; this.height = height;  
    }  
  
    public void runTest(int scrPxs[], int outPxs[]) {  
        int x; int acc;  
        for(x = 0; x < width*height; x++) {  
            acc = (int)((((scrPxs[x] & 0xff) * gMonoMult[0]);  
            acc += (int)((((scrPxs[x] >> 8) & 0xff) * gMonoMult[1]);  
            acc += (int)((((scrPxs[x] >> 16) & 0xff) * gMonoMult[2]);  
            outPxs[x] = (acc) + (acc << 8) + (acc << 16) + (scrPxs[x] << 24);  
        }  
    }  
}
```

# Directive: Target data

**@Target(RENDERScript)**

```
public class GrayScale {  
    private final float gMonoMult[] = ...;
```

**@Map(TO)**

```
    private int width;
```

**@Map(TO)**

```
    private int height;
```

```
    public GrayScale (Activity act, int width, int height) {
```

```
        ...
```

```
    }
```

```
    public void runTest(int scrPxs[], int outPxs[]) {
```

```
        ...
```

```
    }
```

```
}
```

- Constructor method initializes the target context, allocates memory and copy initial values.
- Generate getter or setter methods.
- Generate finalize method that free memory and destroy the target context.

# Directive: Target

**@Target(RENDERScript)**

**public class GrayScale {**

**private final float** gMonoMult[] = ...;

**@Map(TO)**

**private int** width;

**@Map(TO)**

**private int** height;

**public GrayScale** (**Activity act**, **int** width, **int** height) {

...

}

**public void runTest**(**@Map(TO)** **int** scrPxs[], **@Map(FROM)** **int** outPxs[]) {

...

}

}

- Allocate memory and copy values to target context.
- Execute method.
- Copy values from target context and free memory.

# Directive: Parallel

```
@Target(RENDERScript)
```

```
public class GrayScale {  
    private final float gMonoMult[] = ...;  
    @Map(TO)  
    private int width;  
    @Map(TO)  
    private int height;  
    public GrayScale (int width, int height) {  
        ...  
    }  
}
```

- Execute method in parallel.
- Function is executed many times as elements contain the input or output vectors.

```
@Parallel
```

```
public void runTest(@Input int scrPxs[], @Output int outPxs[], @index int x, ...) {  
    int x; int acc;  
    acc = (int)(((scrPxs[x]) & 0xff) * gMonoMult[0]);  
    acc += (int)(((scrPxs[x] >> 8) & 0xff) * gMonoMult[1]);  
    acc += (int)(((scrPxs[x] >> 16) & 0xff) * gMonoMult[2]);  
    outPxs[x] = (acc) + (acc << 8) + (acc << 16) + (scrPxs[x] << 24);  
}  
}
```

ULL

# Directive: Declare

@Target(RENDERSCRIPT)

**public class GrayScale {**

**@Declare**

**private final float** gMonoMult[] = ...;

@Map(TO)

**private int** width;

@Map(TO)

**private int** height;

**public GrayScale** (int width, int height) { ... }

- Fields or methods that are declared only in the target context.

@Parallel

**public void runTest**(@Input int scrPxs[], @Output int outPxs[], @index int x, ...) {

int x; int acc;

acc = (int)((((scrPxs[x]) & 0xff) \* gMonoMult[0]));

acc += (int)((((scrPxs[x] >> 8) & 0xff) \* gMonoMult[1]));

acc += (int)((((scrPxs[x] >> 16) & 0xff) \* gMonoMult[2]));

outPxs[x] = (acc) + (acc << 8) + (acc << 16) + (scrPxs[x] << 24);

}

}

ULL

# COMPUTATIONAL RESULTS

---

# Computational Result

- **Samsung Galaxy SIII**

- Exynos 4 – 4412 (1.4GHz, Quad-core)
- GPU ARM Mali-400/MP4
- 1 GB RAM
- Android 4.1



- **Asus Transformer Prime TF201**

- NVIDIA Tegra 3 (1.4GHz, Quad-core)
- GPU NVIDIA ULP GeForce.
- 1GB RAM
- Android 4.1



- **Asus Nexus 7**

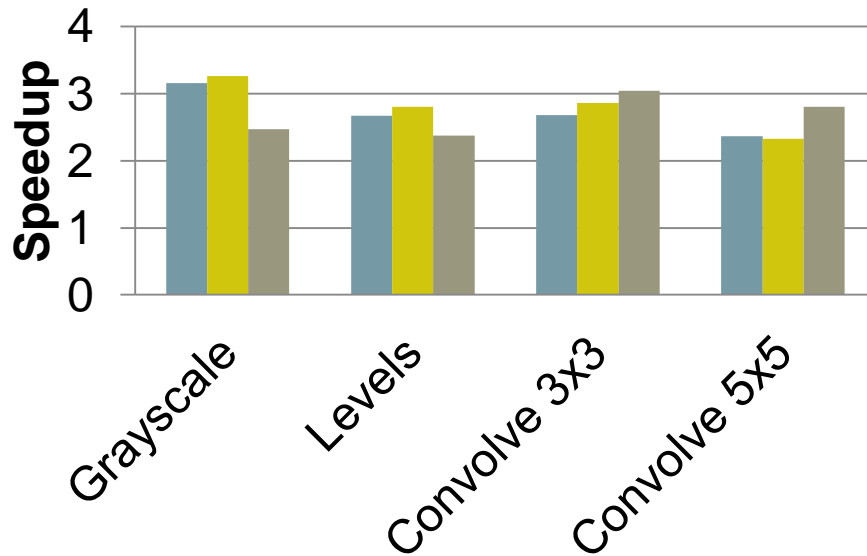
- Qualcomm Snapdragon S4 Pro, (1.5GHz, Quad-core)
- GPU Adreno 320.
- 2GB RAM
- Android 4.3




ULL

# Device comparison

## Native C



## RenderScript

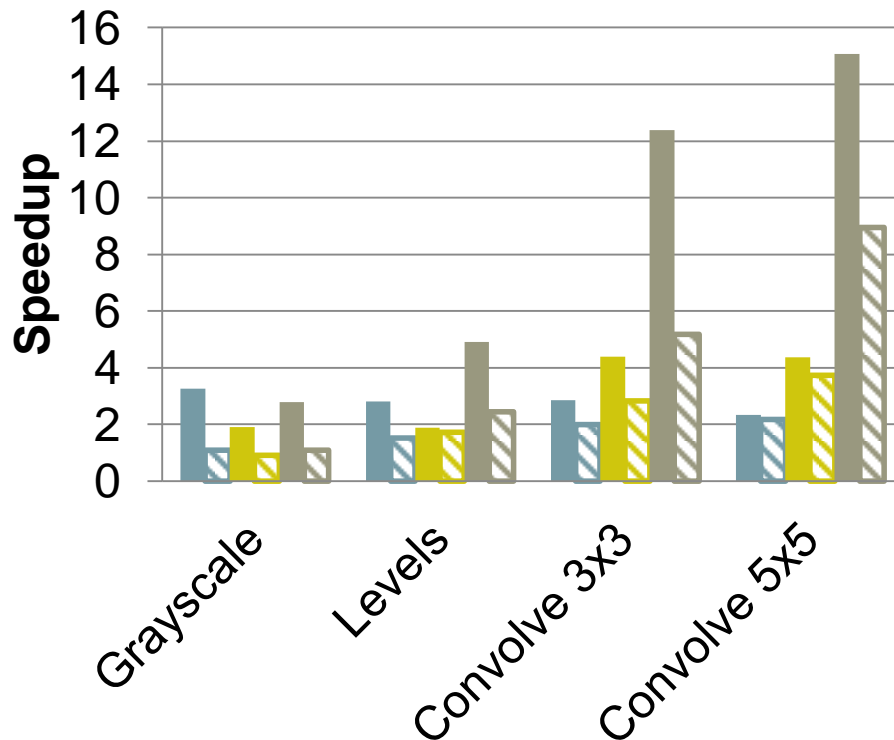


Galaxy SIII  
TF201  
Nexus 7




# Speedup: Galaxy SIII – Asus TF201

## AOSP Benchmark



- Ad-hoc Native C
- Ad-hoc Renderscript Sequential
- Ad-hoc Renderscript Parallel


## General Convolve



- Generated Native C
- Generated Renderscript Sequential
- Generated Renderscript Parallel


# Speedup: Nexus 7

## AOSP Benchmark



- Ad-hoc Native C
- Ad-hoc Renderscript Sequential
- Ad-hoc Renderscript Parallel

## General Convolve



- Generated Native C
- Generated Renderscript Sequential
- Generated Renderscript Parallel

# Timsort

- Hybrid sorting algorithm.
  - Insertion sort.
  - Merge sort.
- Used to sort arrays.
  - Python since version 2.3.
  - Java SE 7.
  - Android.
  - GNU Octave.
- Performance
  - Best case  $O(n)$ .
  - Average case  $O(n \log n)$ .
  - Worst case  $O(n \log n)$ .



# Java TimSort

```
public class TimSort {  
    public void sort(float[] a) {  
        ...  
        BinarySort binarySort = new BinarySort();  
        binarySort.sort(a, loArray, runlenArray, startArray, 0);  
        ...  
        MergeSort merge = new MergeSort();  
        merge.sort(a, loArray, runlenArray);  
    }  
}
```



# Renderscript Binary Sort

```
@Target(RENDERScript)
```

```
public class BinarySort {
```

```
    @Parallel
```

```
    public void sort(float[] a, @Input int loArray[], @Map(TO) int runlenArray[],  
                        @Map(TO) int startArray[], @Index int i) {
```

```
        ...
```

```
        binarySort(a, lo, lo + runLen, lo + start);
```

```
    }
```

```
    @Declare
```

```
    private void binarySort(float[] a, int lo, int hi, int start) {
```

```
        ...
```

```
    }
```

```
}
```



# Native Merge

**@Target(NATIVE)**

```
public class MergeSort {
```

```
    public void sort(float[] a, @Map(TO) int loArray[], @Map(TO) int runlenArray[]) {
```

```
        ...
```

```
        mergeCollapse();
```

```
        ...
```

```
    }
```

**@Declare**

```
    private void mergeCollapse() {
```

```
        ...
```

```
    }
```

```
}
```

# Timsort

- Hybrid sorting algorithm.
  - Insertion sort.
  - Merge sort.
- Used to sort arrays.
  - Python since version 2.3.
  - Java SE 7.
  - Android.
  - GNU Octave.
- Performance
  - Best case  $O(n)$ .
  - Average case  $O(n \log n)$ .
  - Worst case  $O(n \log n)$ .


Java  
Generated Native C  
Generated Renderscript Sequential

Dual Pivot Quicksort


Insertion sort	Merge sort
Java	Native C
RS Parallel	Java
RS Parallel	Native C

# Timsort

Execution time




Speedup



■ Java Quicksort ■ Java ■ Native C ■ Renderscript Sequential



# Timsort



# Conclusion

- The methodology used has been validated on scientific environments.
- We proved that this methodology can be also applied to not scientific environments.
- The tool presented makes easier the development of efficient applications in Android.
- We get efficient code at a low development cost.
- The ad-hoc versions get higher performance but their implementations are more complex.

# Future work

- Adding support for statement annotations.
- Adding new directives and clauses.
- To optimize memory efficiency (supported objects).
- To optimize compute efficiency (vector operation).
- To generate parallel C code.
- To generate parallel Java code.

# THANKS

---

Alejandro Acosta  
aacostad@ull.es

Francisco Almeida  
falmeida@ull.es



Universidad  
de La Laguna

High Performance Computing Group



FEDER-TIN2011-24598